

# Persistence Is Futile

Implementing Delayed Durability  
in SQL Server



**MARK BROADBENT**

Principal Consultant

**sqlcloud**

**SQLCLOUD.CO.UK**

## Contact...



mark.broadbent@sqlcambs.org.uk



@retracement



tenbulls.co.uk

## Likes...



## Guilty pleasures...



**SERVERLESS**

## Badges...

**Microsoft  
CERTIFIED**

Master: SQL Server




## Community...



SQLA



# Agenda



We'll also need to explain In-Memory OLTP in this context!



1

Transactions & Breaking ACID Properties



2

Log Architecture



3

Delayed Durability Implementation On-disk/ IMOLTP



4

Limitations



5

The best part of the presentation... Gin O'Clock

# Requirements

- From SQL 2014 and upward
- On all editions (including Express)
- For SQL on Windows OR SQL on Linux
- In any database recovery mode
  - FULL/ BULK\_LOGGED/ SIMPLE



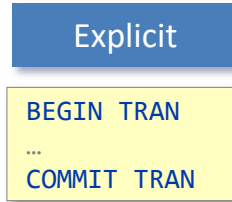




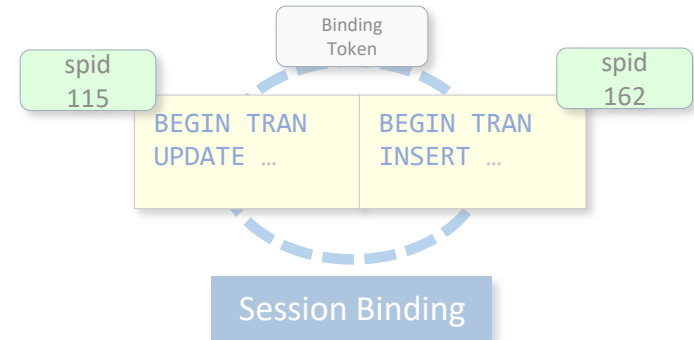
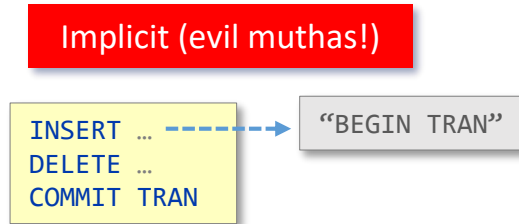
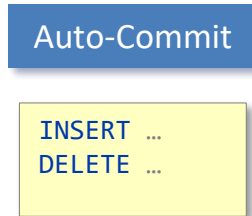
A  
C  
I  
D



# Transactions can be...



(there is also...)



# Transactions

```
BEGIN TRAN transaction_1 WITH MARK 'restorepoint'
  BEGIN TRAN ← "Nested Transaction" *1
    --do something
  COMMIT TRAN
  SAVE TRAN savepoint ← Savepoint
  BEGIN TRAN
    --do something else
  COMMIT
  IF {something_wrong} THEN ROLLBACK TRAN savepoint
  COMMIT ← Outer Transaction
```

*\*1 Not to be confused with Atomic Blocks (we'll discuss these later!)*

# Demo

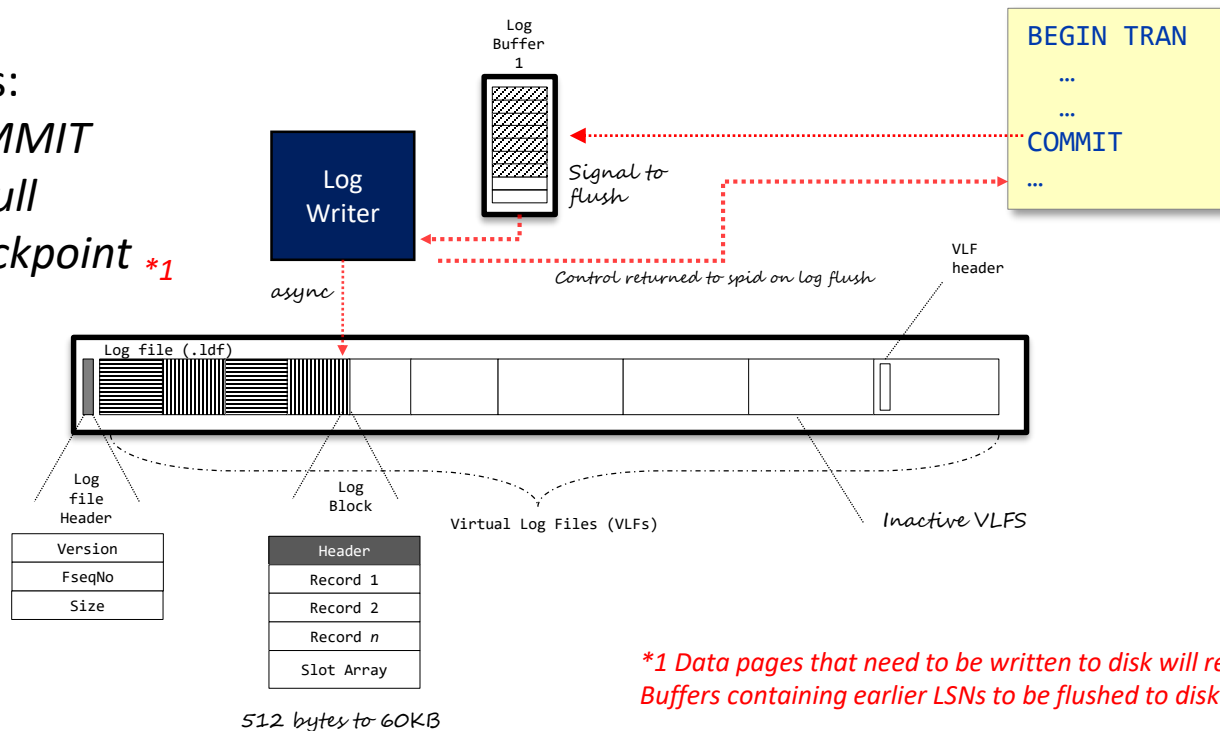
Transactions and broken ACID properties...



# Durable Transactions

Log flushes:

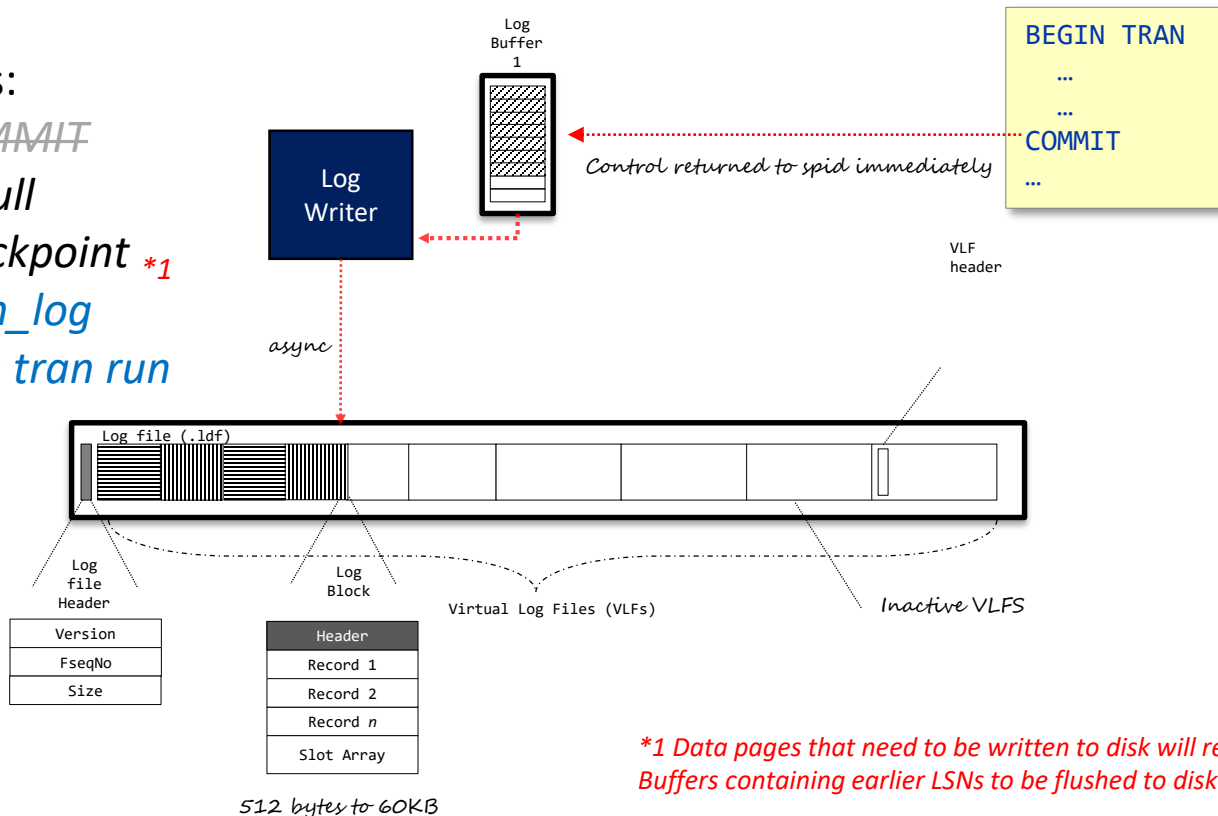
- *On COMMIT*
- *When full*
- *On Checkpoint* \*1



# Delayed Durable Transactions

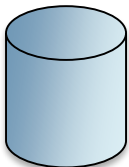
Log flushes:

- ~~On COMMIT~~
- When full
- On Checkpoint \*1
- *sp\_flush\_log*
- Durable tran run



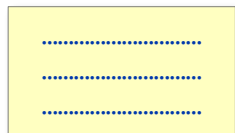
# Delayed Durability Hierachy

## DATABASE



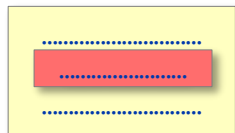
- DISABLED
- ALLOWED
- FORCED

## TRANSACTION (On-disk or IM tables)



- OFF
- ON

## ATOMIC BLOCK (IM tables via Native Compiled Stored Procs)



- OFF
- ON

## Except

- If CDC is enabled
- Is a cross database transaction
- Is a distributed transaction

...BUT delayed durability is  
not guaranteed regardless!

# Demo

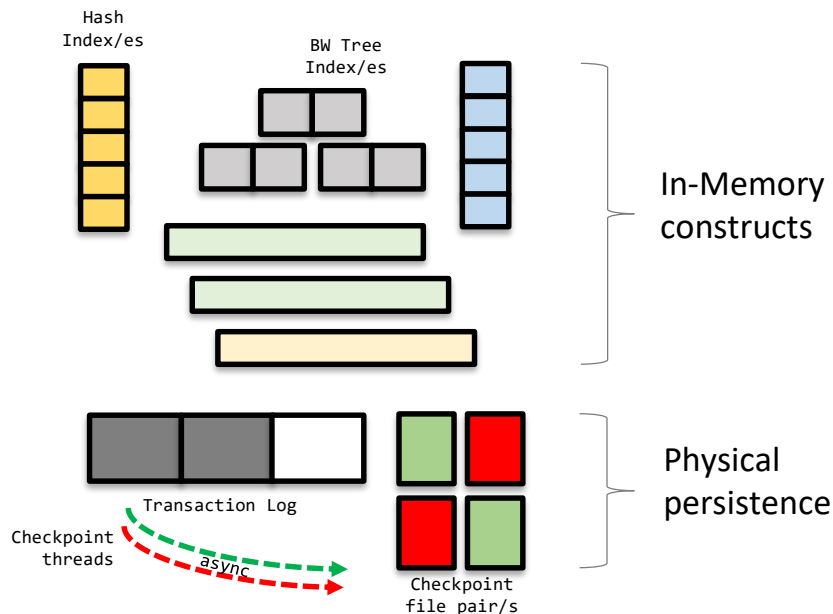
On disk delayed durability

# In-Memory OLTP Overview and Log Improvements

- New concurrency model (as of SQL Server 2014)
- All row and index data in memory, all versioning In-Memory
- Lockless and latchless operation
- Physical persistence when SCHEMA\_AND\_DATA)
  - Thru Checkpoint file pairs + Log

## Log Improvements

- SCHEMA\_ONLY has no log overhead
- Log block “compression” (less log writes)
- No undo record generation
- Indexes not persisted, rebuilt on start-up
  - ...so NO index maintenance logging.



# Demo

In-Memory delayed durability



# Log Waits and Log Performance Monitoring

- LOGBUFFER wait – time taken to create log record in log buffer
- WRITELOG wait – time taken for log buffer to be flushed to logfile
- Use sys.dm\_io\_virtual\_file\_stats to look at file IO, sizes and stalls
- Performance Monitors e.g.
  - Log flushes per second counter of Databases counter set
  - Transactions per second counter of Databases counter set
  - Transactions created per second counter of SQL Server <YYYY> XTP Transactions/ <DB> (instance of object).

# Limitations and Special Cases

- Ignored when
  - CDC enabled tables
  - Cross-database or distributed transactions
- Not supported in Transaction Replication
- Reporting and Failover scenarios *\*1*
  - Sync Read-only replicas
  - High Safety Database Mirrors
  - Failover Clustered Instances
  - Log backups and Log shipping targets
  - Hyper-V Quick Migration.

**Warning!**  
Data loss is very possible in these scenarios!

*\*1 In a nutshell, crash recovery can (and possibly *\*will\**) result in data loss!*

# Summary

- Delayed Durability attempts to make log flushes to disk more efficient
- Is a compromise of performance over durability
- Force on all trans per DB, a specific trans or on atomic block via Natively Compiled SP
- IMOLTP can remove many bottlenecks
  - But IMOLTP already provides some great logging improvements
  - However Delayed Durability \*could\* still remove log buffer waits
- It can result in lost data, even under highly available environments.

# Thank you for listening!

Email: [mark.broadbent@sqlcambs.org.uk](mailto:mark.broadbent@sqlcambs.org.uk)

Twitter: [retracement](#)

Blog: <http://tenbulls.co.uk>

Slideshare: <http://www.slideshare.net/retracement>

Demo: [https://github.com/retracement/Persistence is Futile](https://github.com/retracement/Persistence_is_Futile)