

Persistence Is Futile

Implementing Delayed Durability
in SQL Server



MARK BROADBENT

Principal Consultant

sqlcloud

SQLCLOUD.CO.UK

Contact...



mark.broadbent@sqlcambs.org.uk



@retracement



tenbulls.co.uk

Likes...



Guilty pleasures...



SERVERLESS



Badges...

**Microsoft
CERTIFIED**

Master: SQL Server




Community...



SQLA



Agenda



We'll also need to explain In-Memory OLTP in this context!



1

Transactions & Breaking ACID Properties



2

Log Architecture



3

Delayed Durability Implementation On-disk/ IMOLTP



4

Limitations



5

The best part of the presentation... Gin O'Clock

Requirements



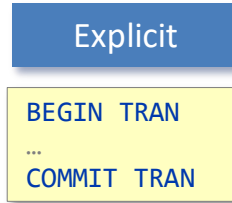
- From SQL 2014 and upward
- On all editions (including Express)
- For SQL on Windows OR SQL on Linux
- In any database recovery mode
 - FULL/ BULK_LOGGED/ SIMPLE



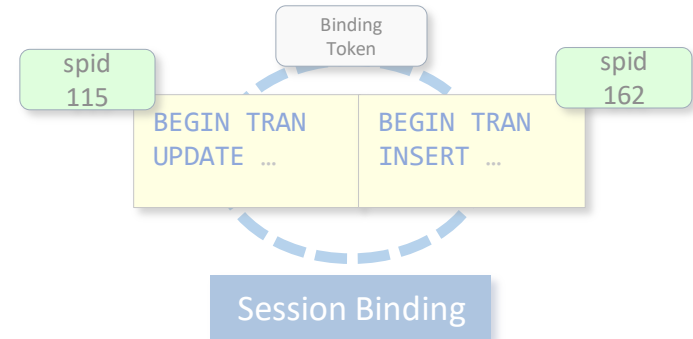
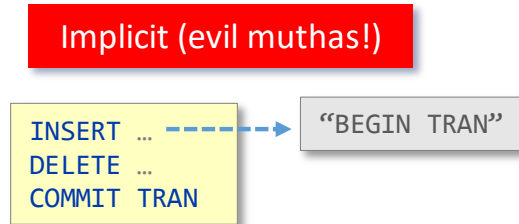
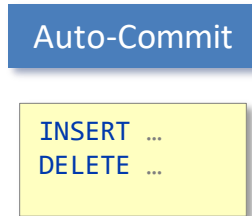
A
C
I
D



Transactions can be...



(there is also...)



Transactions

```
BEGIN TRAN transaction_1 WITH MARK 'restorepoint'
  BEGIN TRAN ← "Nested Transaction" *1
    --do something
  COMMIT TRAN
  SAVE TRAN savepoint ← Savepoint
  BEGIN TRAN
    --do something else
  COMMIT
  IF {something_wrong} THEN ROLLBACK TRAN savepoint
  COMMIT ← Outer Transaction
```

**1 Not to be confused with Atomic Blocks (we'll discuss these later!)*

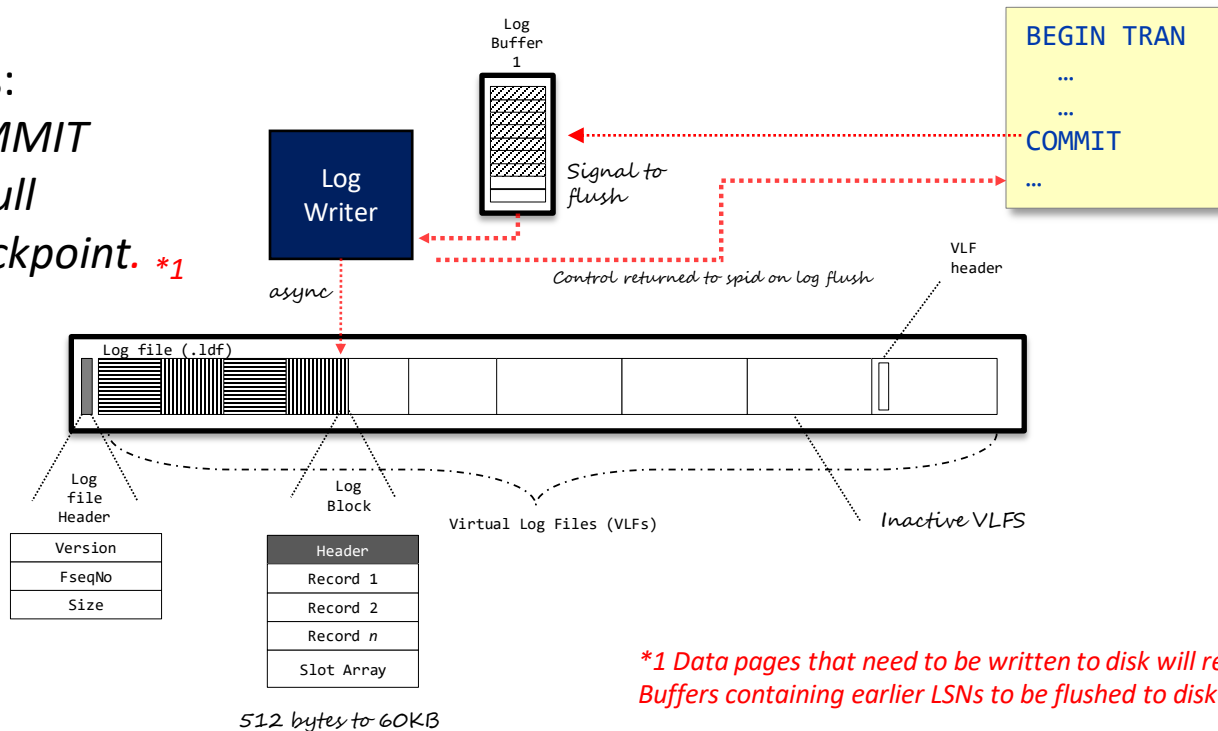
Demo

Transactions and broken ACID properties...

Durable Transactions

Log flushes:


- *On COMMIT*
- *When full*
- *On Checkpoint. *1*



Delayed Durable Transactions

Log flushes:


- ~~On COMMIT~~
- When full
- On Checkpoint. *1
- *sp_flush_log*
- *Durable tran run*



**1 Data pages that need to be written to disk will require Buffers containing earlier LSNs to be flushed to disk*

Delayed Durability Hierachy

DATABASE



- DISABLED
- ALLOWED
- FORCED


TRANSACTION

(On-disk or IM tables)



- OFF
- ON

ATOMIC BLOCK (IM tables via
Native Compiled Stored Procs)



- OFF
- ON

...BUT delayed
durability is
not guaranteed
regardless!

Demo

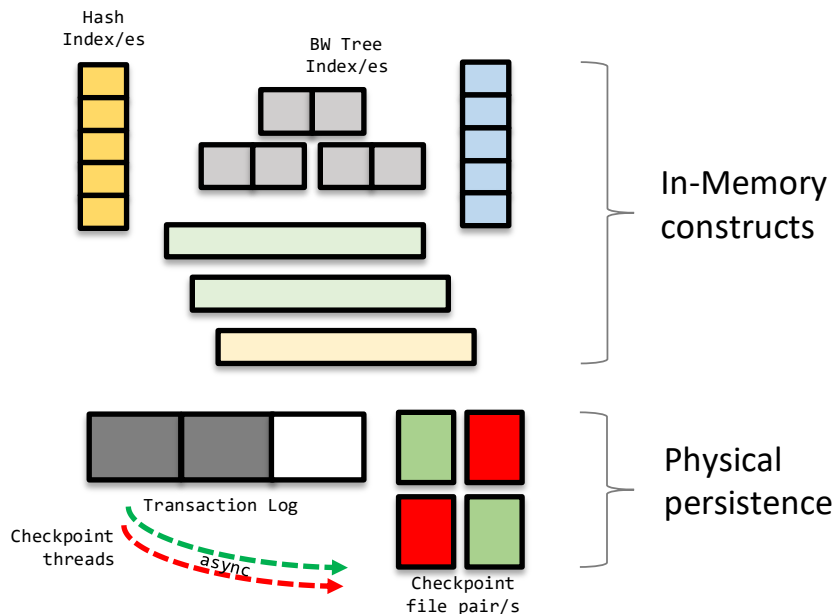
On disk delayed durability

In-Memory OLTP Overview and Log Improvements

- New concurrency model (as of SQL Server 2014)
- All row and index data in memory, all versioning In-Memory
- Lockless and latchless operation
- Physical persistence when SCHEMA_AND_DATA)
 - Thru Checkpoint file pairs + Log

Log Improvements

- SCHEMA_ONLY has no log overhead
- Log block “compression” (less log writes)
- No undo record generation
- Indexes not persisted, rebuilt on start-up
 - ...so NO index maintenance logging.



Demo

In-Memory delayed durability

Log Waits and Log Performance Monitoring

- LOGBUFFER wait – time taken to create log record in log buffer
- WRITELOG wait – time taken for log buffer to be flushed to logfile
- Use sys.dm_io_virtual_file_stats to look at file IO, sizes and stalls
- Performance Monitors e.g.
 - Log flushes per second counter of Databases counter set
 - Transactions per second counter of Databases counter set
 - Transactions created per second counter of SQL Server <YYYY> XTP Transactions/ <DB> (instance of object).

Limitations and Special Cases

- Ignored when
 - CDC enabled tables
 - Cross-database or distributed transactions
- Not supported in Transaction Replication
- Reporting and Failover scenarios *1
 - Sync Read-only replicas
 - High Safety Database Mirrors
 - Failover Clustered Instances
 - Log backups and Log shipping targets
 - Hyper-V Quick Migration.

Warning!
Data loss is very possible in these scenarios!

**1 In a nutshell, crash recovery can (and possibly *will*) result in data loss!*

Summary

- Delayed Durability attempts to make log flushes to disk more efficient
- Is a compromise of performance over durability
- Force on all trans per DB, a specific trans or on atomic block via Natively Compiled SP
- IMOLTP can remove many bottlenecks
 - But IMOLTP already provides some great logging improvements
 - However Delayed Durability *could* still remove log buffer waits
- It can result in lost data, even under highly available environments.

Thank you for listening!

Email: mark.broadbent@sqlcambs.org.uk

Twitter: [retracement](#)

Blog: <http://tenbulls.co.uk>

Slideshare: <http://www.slideshare.net/retracement>

Demo: [https://github.com/retracement/Persistence is Futile](https://github.com/retracement/Persistence_is_Futile)